

CSCI-1200 Data Structures

Final Exam — Practice Problems

*NOTE: The final exam will be cumulative and comprehensive.
This packet contains sample questions from exams from multiple years.*

1 Short Answer [/17]

1.1 Comparing Vectors & Arrays [/5]

The statements below can be used to compare and contrast arrays and vectors. For each statement, specify “**ARRAY**” if it is only true for arrays, “**VECTOR**” if it is only true for vectors, “**BOTH**” if it is true for both types, and “**NEITHER**” if it is true for neither type.

Knows how many elements it contains.

Can be used to store elements of any type.

Prevents access of memory beyond its bounds.

Is dynamically re-sizable.

Can be passed by reference.

1.2 Limited Looping [/3]

True or False There are some algorithms that must be written using a **for** loop and *cannot* be written using a **while** or **do – while** loop.

2 Superhero Division [/14]

In this problem you will add a new operator to the **Superhero** class from lab. Remember that a superhero has a name, a true identity, and a power, but we *cannot* access the true identity of a **Superhero** object from the public interface. Here is the basic **Superhero** class declaration:

```
class Superhero {
public:
    // ACCESSORS
    const string& getName() const { return name; }
    const string& getPower() const { return power; }
    // INPUT STREAM OPERATOR
    friend istream& operator>>(istream &istr, Superhero &hero);
private:
    // REPRESENTATION
    string name;
    string true_identity;
    string power;
};
// OUTPUT STREAM OPERATOR
ostream& operator<<(ostream &ostr, const Superhero &hero);
```

And here is part of the **Superhero** class implementation:

```
ostream& operator<<(ostream &ostr, const Superhero &hero) {
    if (hero.getPower() == "")
        ostr << hero.getName() << " has no power" << endl;
    else
        ostr << "Superhero " << hero.getName() << " has power " << hero.getPower() << endl;
    return ostr;
}
```

Now let's define the `/=` operator on **Superhero**. This operator can be used to defeat a hero by dividing them from their true identity. If an attacker learns a hero's true identity and uses it against them, the superhero loses his power. A superhero must carefully guard his true identity to prevent this attack. If the attacker does not know and just incorrectly guesses the superhero's true identity, this `/=` operation does nothing. For example, suppose `elastigirl` is a **Superhero** object with name equal to "Elastigirl", true identity equal to "Zoe", and power equal to "Flexible". Then the statement:

```
cout << elastigirl;
```

would print this on the screen:

```
Superhero Elastigirl has power Flexible
```

But after executing the statement:

```
elastigirl /= ("Zoe");
```

the output of the variable `elastigirl` would print on the screen as:

```
Elastigirl has no power
```

2.1 Implementation Choices [/5]

Name the three different ways we can implement operator overloading. Which of these three is the most appropriate choice for the /= operator described above? Why?

2.2 /= operator implementation [/9]

Now implement the /= operator. Part of your job is to carefully define the prototype for this function. What should be added or changed in the `superhero.h` class declaration file? And what should be added or changed in the `superhero.cpp` class implementation file? Be specific.

3 Valet Parking Maps [/38]

You have been asked to help with a valet parking system for a big city hotel. The hotel must keep track of all of the cars currently stored in their parking garage and the names of the owners of each car. *Please read through the entire question before working on any of the subproblems.* Here is the simple `Car` class they have created to store the basic information about a car:

```
class Car {
public:
    // CONSTRUCTOR
    Car(const string &m, const string &c) : maker(m), color(c) {}
    // ACCESSORS
    const string& getMaker() const { return maker; }
    const string& getColor() const { return color; }
private:
    // REPRESENTATION
    string maker;
    string color;
};
```

The hotel staff have decided to build their parking valet system using a map between the cars and the owners. This map data structure will allow quick lookup of the owners for all the cars of a particular color and maker (e.g., the owners of all of the silver Hondas in the garage). For example, here is their data structure and how it is initialized to store data about the six cars currently in the garage.

```
map<Car,vector<string> > cars;
cars[Car("Honda","blue")].push_back("Cathy");
cars[Car("Honda","silver")].push_back("Fred");
cars[Car("Audi","silver")].push_back("Dan");
cars[Car("Toyota","green")].push_back("Alice");
cars[Car("Audi","silver")].push_back("Erin");
cars[Car("Honda","silver")].push_back("Bob");
```

The managers also need a function to create a report listing all of the cars in the garage. The statement:

```
print_cars(cars);
```

will result in this report being printed to the screen (`std::cout`):

```
People who drive a silver Audi:
    Dan
    Erin
People who drive a blue Honda:
    Cathy
People who drive a silver Honda:
    Fred
    Bob
People who drive a green Toyota:
    Alice
```

Note how the report is sorted alphabetically by maker, then by car color, and that the owners with similar cars are listed chronologically (the order in which they parked in the garage).

3.1 The `Car` class [/6]

In order for the `Car` class to be used as the first part of a map data structure, what additional non-member function is necessary? Write that function. Carefully specify the function prototype (using `const` & reference as appropriate). Use the example above as a guide.

3.2 Data structure diagram [/10]

Draw a picture of the map data structure stored by the `cars` variable in the example. As much as possible use the conventions from lecture for drawing these pictures. Please be neat when drawing the picture. *Optional: You may also write a few concise sentences to explain your picture.*

3.3 print_cars [/9]

Write the `print_cars` function. Part of your job is to correctly specify the prototype for this function. Be sure to use `const` and pass by reference as appropriate.

3.4 remove_cars [/13]

When guests pick up their cars from the garage, the data structure must be correctly updated to reflect this change. The `remove_car` function returns `true` if the specified car is present in the garage and `false` otherwise.

```
bool success;
success = remove_car(cars, "Erin", "silver", "Audi");
assert (success == true);
success = remove_car(cars, "Cathy", "blue", "Honda");
assert (success == true);
success = remove_car(cars, "Sally", "green", "Toyota");
assert (success == false);
```

After executing the above statements the `cars` data structure will print out like this:

```
People who drive a silver Audi:
  Dan
People who drive a silver Honda:
  Fred
  Bob
People who drive a green Toyota:
  Alice
```

Note that once the only blue Honda stored in the garage has been removed, this color/maker combination is completely removed from the data structure.

Specify the prototype and implement the `remove_car` function.

4 Computational Desert Island [/]

Suppose that a monster is holding you captive on a computational desert island, and has a large file containing double precision numbers that he needs to have sorted. If you write correct code to sort his numbers he will release you and when you return home will be allowed to move on to DSA. If you don't write correct code, he will eventually release you, but only under the condition that you retake CS 1. The stakes indeed are high, but you are quietly confident — you know about the standard library sort function. (Remember, you are supposed to have forgotten all about bubble sort.) The monster startles you by reminding you that this is a computational desert island and because of this the only data structure you have to work with is a queue.

After panicking a bit (or a lot), you calm down and think about the problem. You realize that if you maintain the values in the queue in increasing order, and insert each value into the queue one at a time, then you can solve the rest of the problem easily. Therefore, you must write a function that takes a new double, stored in `x`, and stores it in the queue. Before the function is called, the values in the queue are in increasing order. After the function ends, the values in the queue must also be in increasing order, but the new value must also be among them.

Here is the function prototype:

```
void insert_in_order(double x, queue<double>& q)
```

You may only use the public queue interface (member functions) as specified in lab. You may use a second queue as local variable scratch space or you may try to do it in a single queue (which is a bit harder). Give an “O” estimate of the number of operations required by this function.

5 Operations on Lists [/]

5.1 Reversing a dslist [/]

Write a `dslist<T>` member function called `reverse` that reverses the order of the nodes in the list. The head pointer should point to what was the tail node and the tail pointer should point to what was the head node. All directions of pointers should be reversed. The function prototype is:

```
template <class T> void dslist<T>::reverse();
```

The function must NOT create ANY new nodes.

5.2 Sublists [/]


Write a function to create a new singly-linked list that is a *copy* of a sublist of an existing list. The prototype is:

```
Node<T>* Sublist(Node<T>* head, int low, int high)
```

The `Node` class is:

```
template <class T>
class Node {
public:
    T value;
    Node* next;
};
```

The new list will contain $\text{high} - \text{low} + 1$ nodes, which are copies of the values in the nodes occupying positions low up through and including high of the list pointed to by **head**. The function should return the pointer to the first node in the new list. For example, in the following drawing the original list is shown on top and the new list created by the function when $\text{low} = 2$ and $\text{high} = 4$ is shown below.



A pointer to the first node of this new list should be returned. (In the drawing this would be the value of **nhead**.) You may assume the original list contains at least low nodes. If it contains fewer than high nodes, then stop copying at the end of the original list.

5.3 Splicing into a `cs2list` [/]

Write a `cs2list<T>` member function called `splice` that takes an iterator and a second `cs2list<T>` object and splices the entire contents of the second list between the node pointed to by the iterator and its successor node. The second list must be completely empty afterwards. The function prototype is:

```
template <class T>
void cs2list<T>::splice(iterator itr, cs2list<T>& second);
```

No new nodes should be created by this function AND it should work in $O(1)$ time (i.e. it should be independent of the size of either list).

6 Concurrency and Asynchronous Computing [/3]

Why might a group of dining philosophers starve?

- A) Because it's impossible to eat spaghetti with chopsticks.
- B) Because they are all left-handed.
- C) Because due to a bank error they didn't have enough money in their joint account.
- D) Because they didn't all want to eat at the same time.

7 Garbage Collection [/12]

For each of the real world systems described below, choose the *most appropriate* memory management technique. Each technique should be used exactly once.

- | | |
|-------------------------------------|-----------------------|
| A) Explicit Memory Management (C++) | B) Reference Counting |
| C) Stop & Copy | D) Mark-Sweep |

7.1 Student Registration System [/3]

Must handle the allocation and shuffling of pointers as students register and transfer in and out of classes. Memory usage will not be a deciding factor. Fragmentation of data should be minimized.

7.2 Playing Chess [/3]

Implementation of a tree-based algorithm for searching the game space. Remember that a *tree* is a *graph* with no cycles.

7.3 Webserver [/3]

A collection of infrequently changing interconnected webpages. Any memory usage overhead should be low. Pauses in service are tolerable.

7.4 Hand Held Game (e.g., GameBoy or PSP, etc.) [/3]

Performance critical application with extremely limited memory resources.

8 Short Answer [/22]

8.1 Garbage Identification [/7]

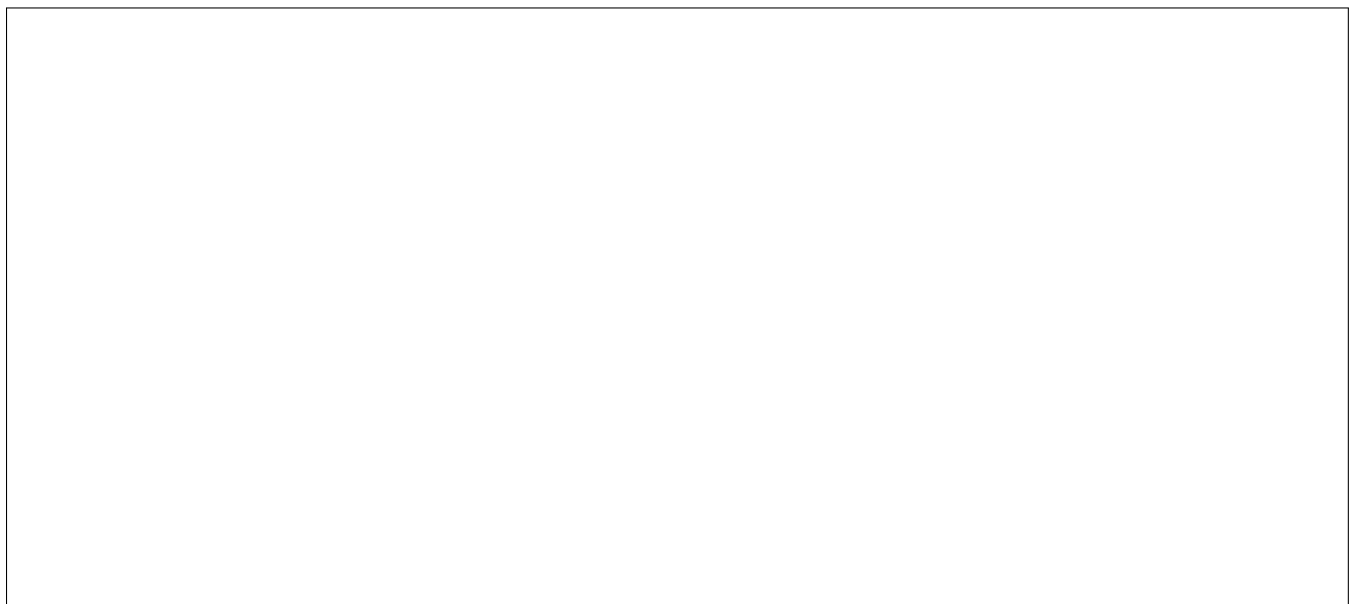
To which address in the memory below should the root variable point so that exactly 2 cells are garbage? Draw a *box and pointer diagram* to justify your answer and state which 2 cells are garbage.

address	100	101	102	103	104	105	106	107
value	a	b	c	d	e	f	g	h
left	106	106	107	100	102	101	0	0
right	103	105	105	0	105	101	0	101



8.2 Stop and Copy Garbage Collection [/4]

What is the purpose of the *forwarding address* in Stop and Copy garbage collection? What will go wrong if you neglect to record this value? Write 2 or 3 concise and well-written sentences.



8.3 Concurrency and Asynchronous Computing [/4]

When programming with multiple threads or processes, the correct use of mutexes (locks) and condition variables will ensure that:

- A) The program always returns the exact same answer.
- B) The program returns an answer that was not possible if the program ran sequentially.
- C) The entire program is atomic.
- D) Each student in a large class will be able to successfully copy a complete set of lecture notes (with no repetitions), even if there are multiple professors.
- E) Deadlock will be avoided if there are multiple mutexes, but may still happen in systems with a single lock.

8.4 Perfect Hashing for Image Compression [/7]

For the last homework, you implemented a compression scheme for 2D images. What are the drawbacks of using this format as the underlying representation for an image editing program? What types of edits to the image are simple? What types of edits will be comparatively inefficient to process? Write 3-4 *concise and well-written* sentences.

9 Data Structures [/18]

Indicate by letter the data structure(s) that have each characteristic listed below.

- | | | | |
|-------------------|---------------|-----------------|--------|
| A) vector | B) list | C) map | D) set |
| E) priority queue | F) hash table | G) leftist heap | |

allows efficient (sublinear) removal of the first and last elements (or the minimum and maximum elements)

uses an array or vector as the underlying representation

uses a network of nodes connected by pointers as the underlying representation

the underlying data structure must be “balanced” or well-distributed to achieve the targeted performance

requires definition of `operator<` or `operator>`

entries cannot be modified after they are inserted (requires re-insertion or re-processing of position)

duplicates are not allowed

allows sublinear merging of two of instances of this data structure

10 Order Notation [/16]

Match the order notation with each fragment of code. Two of the letters will not be used.

- | | | | |
|-------------|----------------|------------------|------------------|
| A) $O(n)$ | B) $O(1)$ | C) $O(n^n)$ | D) $O(n^2)$ |
| E) $O(2^n)$ | F) $O(\log n)$ | G) $O(n \log n)$ | H) $O(\sqrt{n})$ |

```
vector<int> my_vector;
// my_vector is initialized with n entries
// do not include initialization in performance analysis
for (int i = 0; i < n; i++) {
    my_vector.erase(my_vector.begin());
}
```

```
map<string,int> my_map;
// my_map is initialized with n entries
// do not include initialization in performance analysis
my_map.find("hello");
```

```
int foo(int n) {
    if (n == 1 || n == 0) return 1;
    return foo(n-1) + foo(n-2);
}
```

```
int k = 0;
for (int i = 0; i < sqrt(n); i++) {
    for (int j = 0; j < sqrt(n); j++) {
        k += i*j;
    }
}
```

```
set<string> my_set;
for (int i = 0; i < n; i++) {
    string s;
    cin >> s;
    my_set.insert(s);
}
```

```
float* my_array = new float[n];
// do not include memory allocation in performance analysis
my_array[n/2] = sqrt(n);
```


11 Office Demolition [/31]

In this problem we will explore a simple class to manage the assignment of people to offices and desks. Each `Office` object stores its name, the number of desks it can hold, and the names of the people assigned to those desks. An office also stores a *reference* to a master queue of all the people who still need to be assigned to desks. When an office is constructed, people are assigned to the office from the front of this master queue. When an office is demolished, the people who were assigned to that office should be added to the end of the queue while they wait for a new office assignment. Here is the *partial* declaration of the `Office` class:

```
class Office {
public:
    Office(const string& name, int num_desks, queue<string> &unassigned);
    friend ostream& operator<<(ostream &ostr, const Office &office);
private:
    // representation
    string _name;
    int _num_desks;
    string* _desks;
    queue<string>& _unassigned; // a reference to the master queue
};
```

In the example below we create the master queue of people who need to be assigned to desks in offices, and create and delete several `Office` objects:

```
queue<string> unassigned;
unassigned.push("Alice");
unassigned.push("Bob");
unassigned.push("Cathy");
unassigned.push("Dan");
unassigned.push("Erin");
unassigned.push("Fred");
unassigned.push("Ginny");

Office *red = new Office("red", 4, unassigned);
Office *green = new Office("green", 2, unassigned);
cout << *red << *green;

delete red;
cout << "After deleting the red office, "
    << unassigned.size() << " people are waiting for desks." << endl;

Office *blue = new Office("blue", 3, unassigned);
cout << *blue;

cout << "Before deleting the blue & green offices, "
    << unassigned.size() << " people are waiting for desks." << endl;
delete green;
delete blue;
cout << "After deleting all of the offices, "
    << unassigned.size() << " people are waiting for desks." << endl;
```

Here is the desired output from this example:

```
The red office has 4 desks:
  desk[0] = Alice
  desk[1] = Bob
  desk[2] = Cathy
  desk[3] = Dan
The green office has 2 desks:
  desk[0] = Erin
  desk[1] = Fred
After deleting the red office, 5 people are waiting for desks.
The blue office has 3 desks:
  desk[0] = Ginny
  desk[1] = Alice
  desk[2] = Bob
Before deleting the blue & green offices, 2 people are waiting for desks.
After deleting all of the offices, 7 people are waiting for desks.
```

Here is the implementation of the constructor, as it appears in the `office.cpp` file:

```
Office::Office(const string& name, int num_desks, queue<string> &unassigned)
: _name(name), _num_desks(num_desks), _unassigned(unassigned) {
  _desks = new string[_num_desks]; // allocate the desk space
  for (int i = 0; i < _num_desks; i++) {
    if (_unassigned.size() > 0) { // assign from the master queue
      _desks[i] = _unassigned.front();
      _unassigned.pop();
    } else { // if there are no unassigned people, leave the desk empty
      _desks[i] = "";
    }
  }
}
```

11.1 Classes and Memory Allocation [/10]

Anytime you write a new class, especially those with *dynamically allocated memory*, it is very important to consider the member functions that the compiler will automatically generate and determine if this default behavior is appropriate. List these 4 important functions by their generic names, *AND* write their prototypes as they would appear within the `Office` class declaration.

11.2 Declaring a Destructor [/3]

The `Office` class is incomplete and requires implementation of a custom destructor so that people assigned to demolished offices are returned to the master queue and memory is deallocated as appropriate to avoid memory leaks. What line needs to be added to the header file to declare the destructor? Be precise with syntax. Where should this line be added: within the `public`, `protected`, or `private` interface?

11.3 Implementing a Destructor [/12]

Implement the destructor, as it would appear in the `office.cpp` file.

11.4 Operator Overloading [/6]

Here is the implementation of the << stream operator as it appears within the `office.cpp` file:

```
ostream& operator<<(ostream &ostr, const Office &o) {
    ostr << "The " << o._name << " office has "
        << o._num_desks << " desks:" << endl;
    for (int i = 0; i < o._num_desks; i++) {
        ostr << "  desk[" << i << "] = " << o._desks[i] << endl;
    }
    return ostr;
}
```

There are three different ways to overload an operator: as a non-member function, as a member function, and as a friend function. Which method was selected for the `Office` object << stream operator? What are the reasons for this choice? Discuss why the other two methods are inappropriate or undesirable. Write 3 or 4 concise and thoughtful sentences.

12 Dynamically-Allocated Arrays [/17]

Write a function that takes an STL list of integers, finds the even numbers, and places them in a dynamically-allocated array. Only the space needed for the even numbers should be allocated, and no containers other than the given list and the newly-created array may be used. As an example, given a list containing the values:

```
3 10 -1 5 6 9 13 14
```



the function should allocate an array of size 3 and store the values 10, 6 and 14 in it. It should return, via arguments, both the pointer to the start of the array and the number of values stored. No subscripting may be used — not even `*(a+i)` in place of `a[i]`. Here is the function prototype:

```
void even_array(const list<int>& b, int* & a, int& n);
```

13 Ternary Tree Recursion [/17]

A *ternary tree* is similar to a binary tree except that each node has at most 3 children. Write a *recursive* function named `EqualsChildrenSum` that takes one argument, a pointer to the root of a ternary tree, and returns true if the value at each non-leaf node is the sum of the values of all of its children and false otherwise. In the examples below, the tree on the left will return true and the tree on the right will return false.

```
class Node {  
public:  
    int value;  
    Node* left;  
    Node* middle;  
    Node* right;  
};
```



14 Priority Queues [/16]

```
template <class T> class priority_queue {
public:
    // CONSTRUCTOR
    priority_queue() {}
    // ACCESSORS
    int size() { return m_heap.size(); }
    bool empty() { return m_heap.empty(); }
    const T& top() const { assert(!m_heap.empty()); return m_heap[0]; }
    // MODIFIERS
    void push(const T& entry) {
        m_heap.push_back(entry);
        this->percolate_up(int(m_heap.size()-1));
    }
    void pop() { // find and remove the element with the smallest value
        assert(!m_heap.empty());
        m_heap[0] = m_heap.back();
        m_heap.pop_back();
        this->percolate_down(0);
    }
    void pop_max() { /* YOU WILL IMPLEMENT THIS FUNCTION */ }

private:
    // HELPER FUNCTIONS
    void percolate_up(int i) {
        T value = m_heap[i];
        while (i > 0) {
            int parent = (i-1)/2;
            if (value >= m_heap[parent]) break; // done
            m_heap[i] = m_heap[parent];
            i = parent;
        }
        m_heap[i] = value;
    }
    void percolate_down(int i) {
        T value = m_heap[i];
        int last_non_leaf = int(m_heap.size()-1)/2;
        while (i <= last_non_leaf) {
            int child = 2*i+1, rchild = 2*i+2;
            if (rchild < m_heap.size() && m_heap[child] > m_heap[rchild])
                child = rchild;
            if (m_heap[child] >= value) break; // found right location
            m_heap[i] = m_heap[child];
            i = child;
        }
        m_heap[i] = value;
    }

    // REPRESENTATION
    vector<T> m_heap;
};
```

14.1 Implementing `pop_max` [/12]

Write the new priority queue member function named `pop_max` that finds and removes from the queue the element with the *largest* value. Carefully think about the efficiency of your implementation. Remember that a standard priority queue stores the smallest value element at the root.

14.2 Analysis [/4]

If there are n elements in the priority queue, how many elements are visited by the `pop_max` function in the worst case? What is the order notation for the running time of this function?

15 Inheritance & Polymorphism [/10]

What is the output of the following program?

```
class A {
public:
    virtual void f() { cout << "A::f\n"; }
    void g() { cout << "A::g\n"; }
};

class B : public A {
public:
    void g() { cout << "B::g\n"; }
};

class C : public B {
public:
    void f() { cout << "C::f\n"; }
    void g() { cout << "C::g\n"; }
};

int main() {
    A* a[3];
    a[0] = new A();
    a[1] = new B();
    a[2] = new C();

    for (int i = 0; i < 3; i++) {
        cout << i << endl;
        a[i]->f();
        B* b = dynamic_cast<B*>(a[i]);
        if (b) b->g();
    }
}
```

16 Types & Values [/15]

For the *last expression* in each fragment of code below, give the *type* (`int`, `vector<double>`, `Foo*`, etc.) and the *value*. If the value is a legal address in memory, write “**memory address**”. If the value hasn’t been properly initialized, write “**uninitialized**”. If there is an error in the code, write “**error**”. You may want to draw a picture to help you answer each question, but credit will only be given for what you’ve written in the boxes.

```
double a = 5.2;
double b = 7.5;
a+b
```

Type:

Value:

```
int *d;
int e[7] = { 15, 6, -7, 19, -1, 3, 22 };
d = e + e[5];
*d
```

Type:

Value:

```
bool *f = new bool;
*f = false;
f
```

Type:

Value:

```
int g = 10;
int *h = new int[g];
h[0]
```

Type:

Value:

```
map<string, int> m;
m.insert(make_pair(string("bob"),5551111));
m.insert(make_pair(string("dave"),5552222));
m.insert(make_pair(string("alice"),5553333));
m.insert(make_pair(string("chris"),5554444));
(++m.find("bob"))->second
```

Type:

Value: